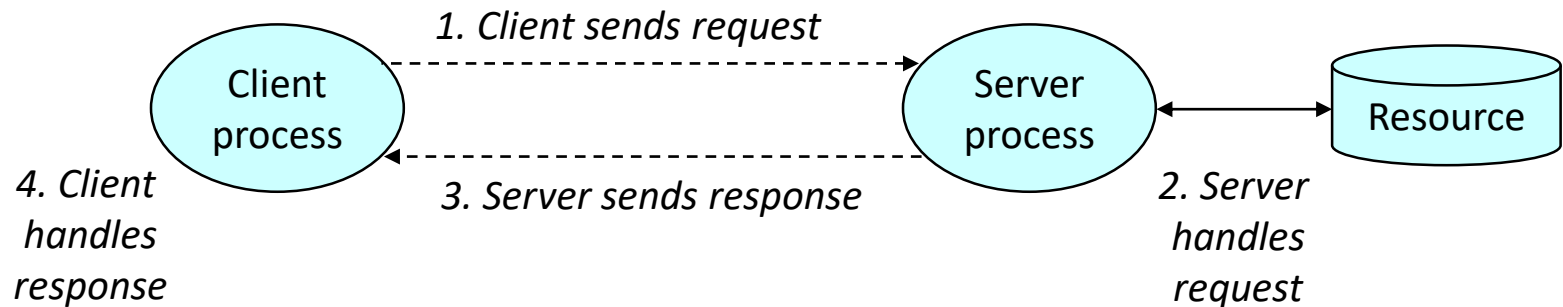


Programmer's view of the Internet

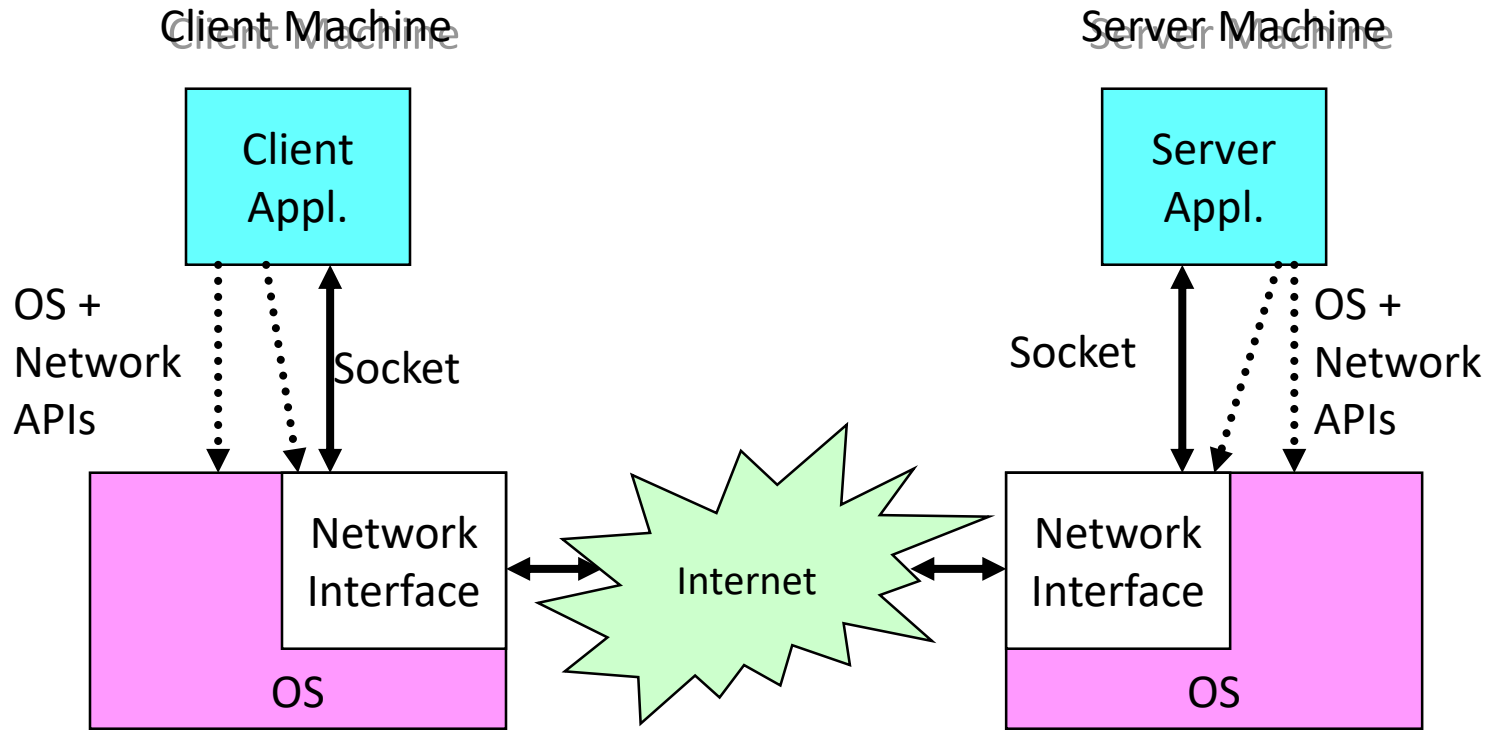
A Client-Server Exchange

- A *server* process and one or more *client* processes
- Server manages some *resource*.
- Server provides *service* by manipulating resource for clients.



Note: clients and servers are processes running on hosts (can be the same or different hosts).

Network Applications



Access to Network via Program Interface

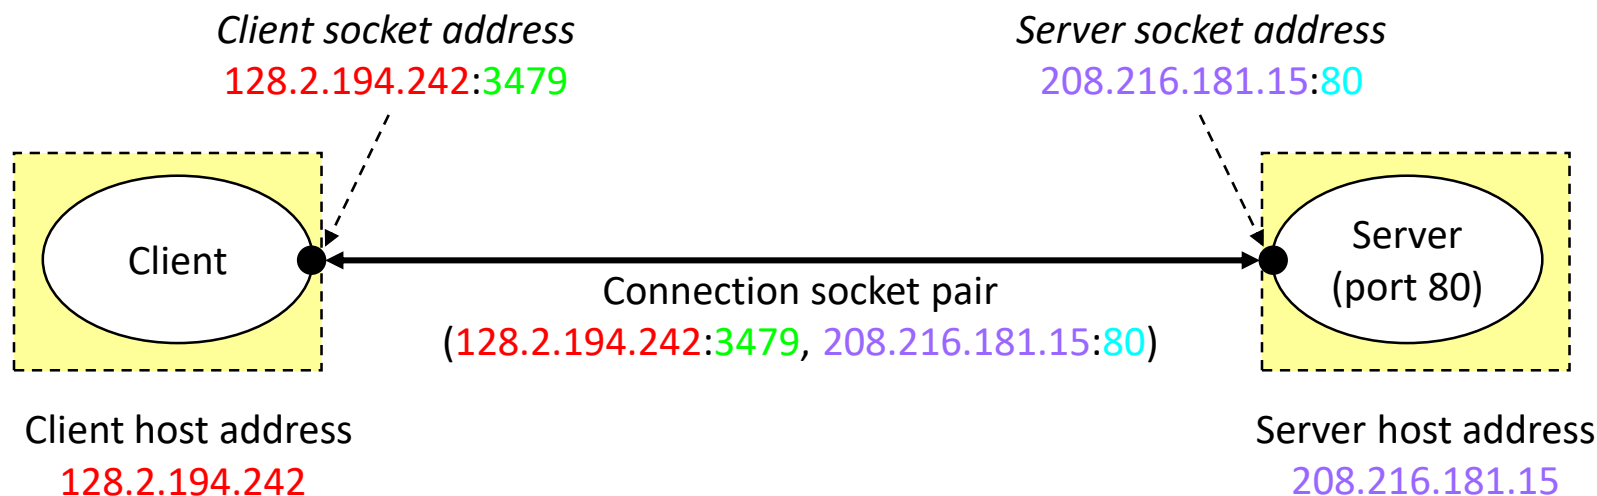
- Sockets make network I/O look like files
- Call system functions to control and communicate
- Network code handles issues of routing, segmentation.

Internet Connections (TCP/IP)

Two common paradigms for clients and servers communication

- Datagram Socket(UDP protocol SOCK_DGRAM)
- Stream Socket (TCP protocol, SOCK_STREAM)

Connections are point-to-point, full-duplex (2-way communication), and reliable.



Note: 3479 is an ephemeral port allocated by the kernel

Note: 80 is a well-known port associated with Web servers

Clients

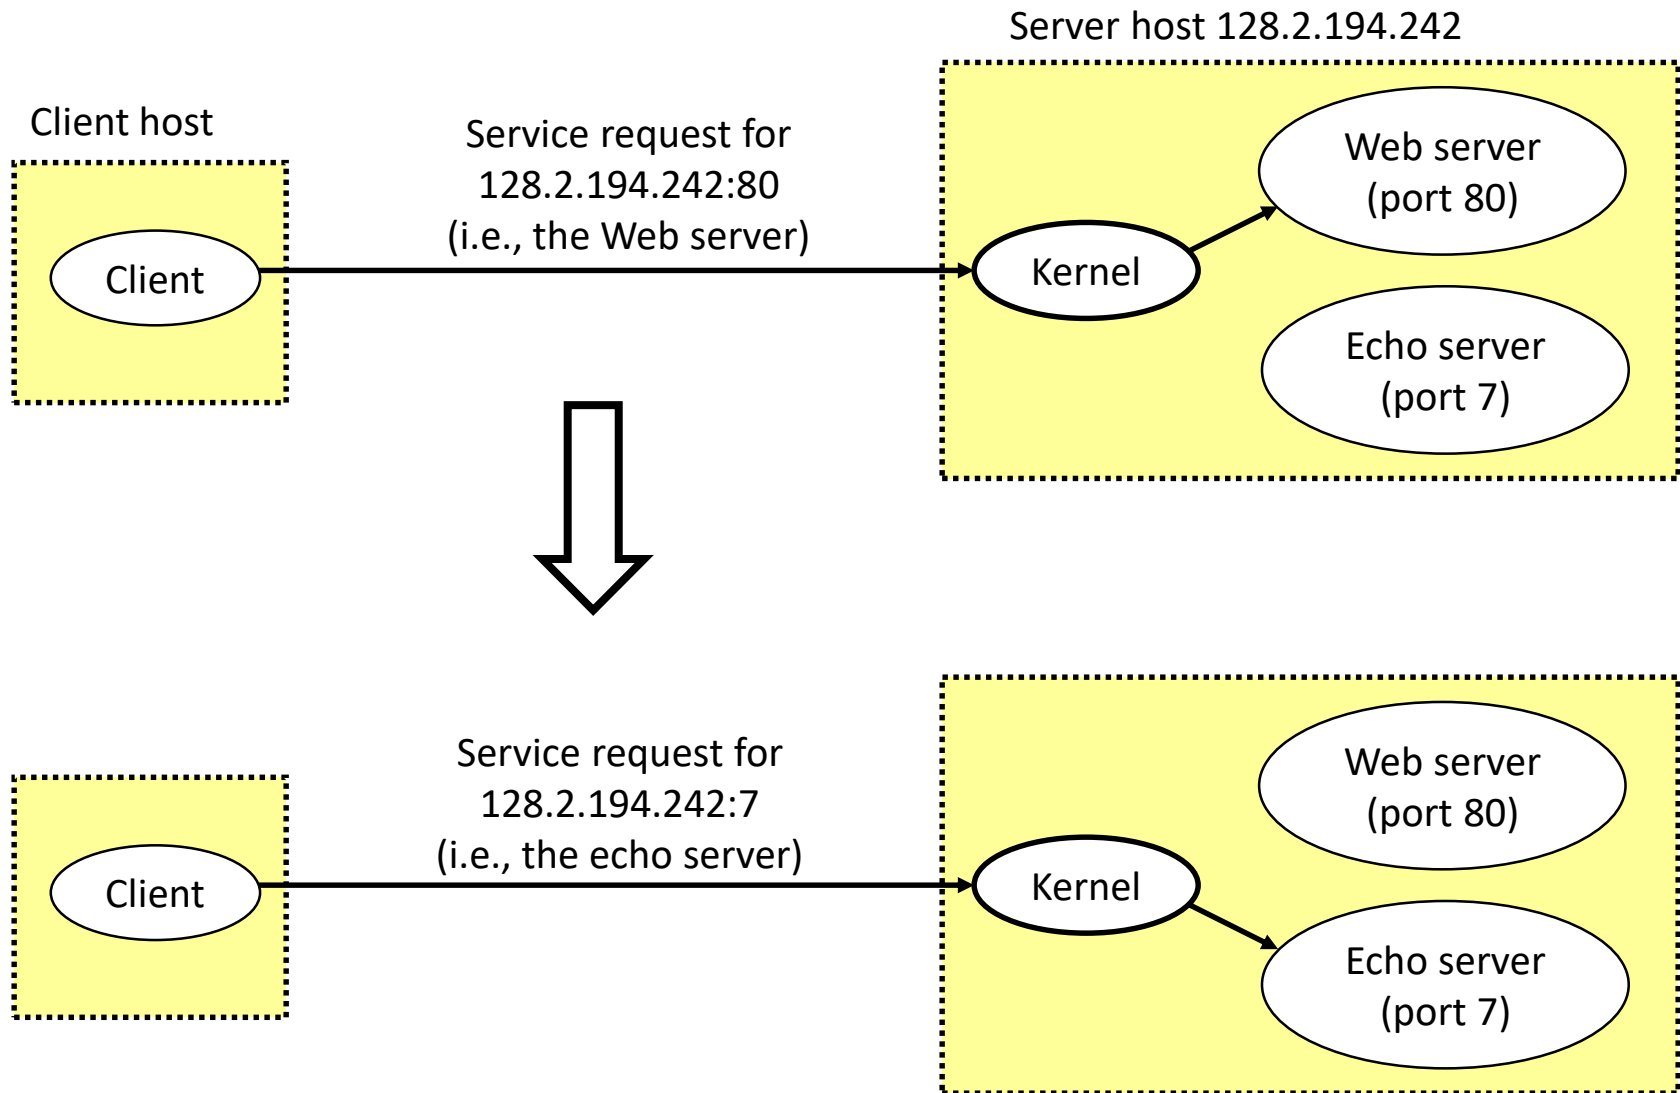
Examples of client programs

- Web browsers, `ftp`, `telnet`, `ssh`

How does a client find the server?

- The IP address in the server socket address identifies the host (*more precisely, an adaptor on the host*)
- The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
- Examples of well known ports
 - Port 7: Echo server
 - Port 23: Telnet server
 - Port 25: Mail server
 - Port 80: Web server

Using Ports to Identify Services



Servers

- Servers are long-running processes (daemons).
 - Created at boot-time (typically) by the init process (process 1)
 - Run continuously until the machine is turned off.
- Each server waits for requests to arrive on a well-known port associated with a particular service.
 - Port 7: echo server
 - Port 23: telnet server
 - Port 25: mail server
 - Port 80: HTTP server
- A machine that runs a server process is also often referred to as a “server.”

See `/etc/services` for a comprehensive list of the services available on a Linux machine.

Sockets Introduction

Socket Address Structures

- An IPv4 socket address structure, commonly called an "Internet socket address structure,"
- It is named `sockaddr_in` and is defined by including the `<netinet/in.h>` header

```
struct in_addr {
    in_addr_t    s_addr;          /* 32-bit IPv4 address */
                                   /* network byte ordered */
};

struct sockaddr_in {
    uint8_t       sin_len;        /* length of structure (16) */
    sa_family_t   sin_family;     /* AF_INET */
    in_port_t     sin_port;       /* 16-bit TCP or UDP port number */
                                   /* network byte ordered */
    struct in_addr sin_addr;      /* 32-bit IPv4 address */
                                   /* network byte ordered */
    char          sin_zero[8];    /* unused */
};
```

Datatypes required by the POSIX specification

Datatype	Description	Header
int8_t	Signed 8-bit integer	<sys/types.h>
uint8_t	Unsigned 8-bit integer	<sys/types.h>
int16_t	Signed 16-bit integer	<sys/types.h>
uint16_t	Unsigned 16-bit integer	<sys/types.h>
int32_t	Signed 32-bit integer	<sys/types.h>
uint32_t	Unsigned 32-bit integer	<sys/types.h>
sa_family_t	Address family of socket address structure	<sys/socket.h>
socklen_t	Length of socket address structure, normally uint32_t	<sys/socket.h>
in_addr_t	IPv4 address, normally uint32_t	<netinet/in.h>
in_port_t	TCP or UDP port, normally uint16_t	<netinet/in.h>

Generic Socket Address Structure

- How to declare the type of pointer that may accept any type of pointer?
- With ANSI C, the solution is simple: `void *` is the generic pointer type
- The socket functions when take an argument it is defining a variable of *generic* socket address structure define in the `<sys/socket.h>`

```
struct sockaddr {  
    uint8_t      sa_len;  
    sa_family_t  sa_family;    /* address family: AF_XXX value */  
    char         sa_data[14]; /* protocol-specific address */  
};
```

IPv6 Socket Address Structure

- The IPv6 socket address is defined by including the `<netinet/in.h>` header

```
struct in6_addr {
    uint8_t  s6_addr[16];          /* 128-bit IPv6 address */
                                    /* network byte ordered */
};

#define SIN6_LEN      /* required for compile-time tests */

struct sockaddr_in6 {
    uint8_t      sin6_len;          /* length of this struct (28) */
    sa_family_t  sin6_family;      /* AF_INET6 */
    in_port_t    sin6_port;        /* transport layer port# */
                                    /* network byte ordered */
    uint32_t     sin6_flowinfo;    /* flow information, undefined */
    struct in6_addr sin6_addr;      /* IPv6 address */
                                    /* network byte ordered */
    uint32_t     sin6_scope_id;    /* set of interfaces for a scope */
};
```

New Generic Socket Address Structure

- A new generic socket address structure was defined as part of the IPv6 sockets API
- This is to overcome some of the shortcomings of the existing `struct sockaddr`
- The new `struct sockaddr_storage` is large enough to hold any socket address type supported by the system


```
struct sockaddr_storage {  
    uint8_t      ss_len;          /* length of this struct (implementation dependent)  
    sa_family_t  ss_family;      /* address family: AF_XXX value */  
    /* implementation-dependent elements to provide:  
    * a) alignment sufficient to fulfill the alignment requirements of  
    *    all socket address types that the system supports.  
    * b) enough storage to hold any type of socket address that the  
    *    system supports.  
    */  
};
```

Value-Result Arguments

- We mentioned that when a socket address structure is passed to any socket function, it is always passed by reference
 - That is, a pointer to the structure is passed
- The length of the structure is also passed as an argument
 - But the way in which the length is passed depends on which direction the structure is being passed:
 - from the process to the kernel
 - bind, connect, and sendto, pass a socket address structure from the process to the kernel
 - or vice versa
 - accept, recvfrom, getsockname, and getpeername, pass a socket address structure from the kernel to the process

Socket address structure passed from process to kernel


```
#define SA      struct sockaddr  
  
struct sockaddr_in serv;  
  
/* fill in serv{} */  
connect (sockfd, (SA *) &serv, sizeof(serv));
```



Socket address structure passed from kernel to process

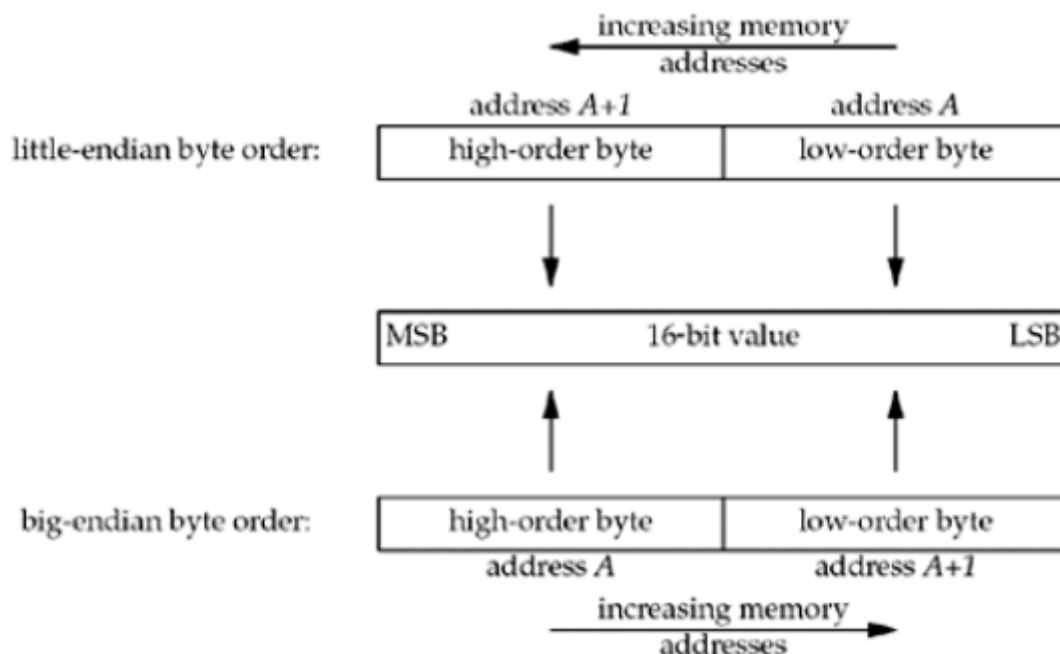
```
struct sockaddr_un cli;    /* Unix domain */
socklen_t len;

len = sizeof(cli);         /* len is a value */
getpeername(unixfd, (SA *) &cli, &len);
/* len may have changed */
```



Byte Ordering Functions

- Consider a 16-bit integer that is made up of 2 bytes
- There are two ways to store the two bytes in memory:
 - with the low-order byte at the starting address, known as *little-endian* byteorder
 - with the high-order byte at the starting address, known as *big-endian* byte order



Byte Manipulation Functions

- There are two groups of functions that operate on multibyte fields, without interpreting the data, and without assuming that the data is a null-terminated C string
- We need these types of functions when dealing with socket address structures because we need to manipulate fields such as IP addresses
- The first group of functions, whose names begin with b (for byte), are from 4.2BSD and are still provided by almost any system that supports the socket functions
- The second group of functions, whose names begin with mem (for memory), are from the ANSI C standard and are provided with any system that supports an ANSI C library

Byte Manipulation Functions

```
#include <strings.h>
```

```
void bzero(void *dest,size_tnbytes);
```

```
void bcopy(const void *src,void *dest,size_tnbytes);
```

```
int bcmp(const void *ptr1,const void *ptr2,size_tnbytes);
```

Returns: 0 if equal, nonzero if unequal

```
#include <string.h>
```

```
void *memset(void *dest,intc,size_tlen);
```

```
void *memcpy(void *dest,const void *src,size_tnbytes);
```

```
int memcmp(const void *ptr1,const void *ptr2,size_tnbytes);
```

Returns: 0 if equal, <0 or >0 if unequal (see text)

inet_aton, inet_addr, and inet_ntoa Functions

- inet_aton, inet_ntoa, and inet_addr convert an IPv4 address from a dotted-decimal string (e.g., "206.168.112.96") to its 32-bit network byte ordered binary value

<code>#include <arpa/inet.h></code>	
<code>int inet_aton(const char *strptr, struct in_addr *addrptr);</code>	
	Returns: 1 if string was valid, 0 on error
<code>in_addr_t inet_addr(const char *strptr);</code>	
	Returns: 32-bit binary network byte ordered IPv4 address; <code>INADDR_NONE</code> if error
<code>char *inet_ntoa(struct in_addr inaddr);</code>	
	Returns: pointer to dotted-decimal string

`int inet_aton(const char *strptr, struct in_addr *addrptr);`

- `inet_aton`, converts the C character string pointed to by *strptr* into its 32-bit binary network byte ordered value, which is stored through the pointer *addrptr*
- If successful, 1 is returned; otherwise, 0 is returned

`in_addr_t inet_addr(const char *strptr);`

- `inet_addr` does the same conversion, returning the 32-bit binary network byte ordered value as the return value

`char *inet_ntoa(struct in_addr inaddr);`

- The `inet_ntoa` function converts a 32-bit binary network byte ordered IPv4 address into its corresponding dotted-decimal string

inet_pton and inet_ntop Functions

- These two functions work with both IPv4 and IPv6 addresses
- The letters "p" and "n" stand for *presentation* (ASCII string) and *numeric* (binary value)

```
#include <arpa/inet.h>
```

```
int inet_pton(int family, const char *strptr, void *addrptr);
```

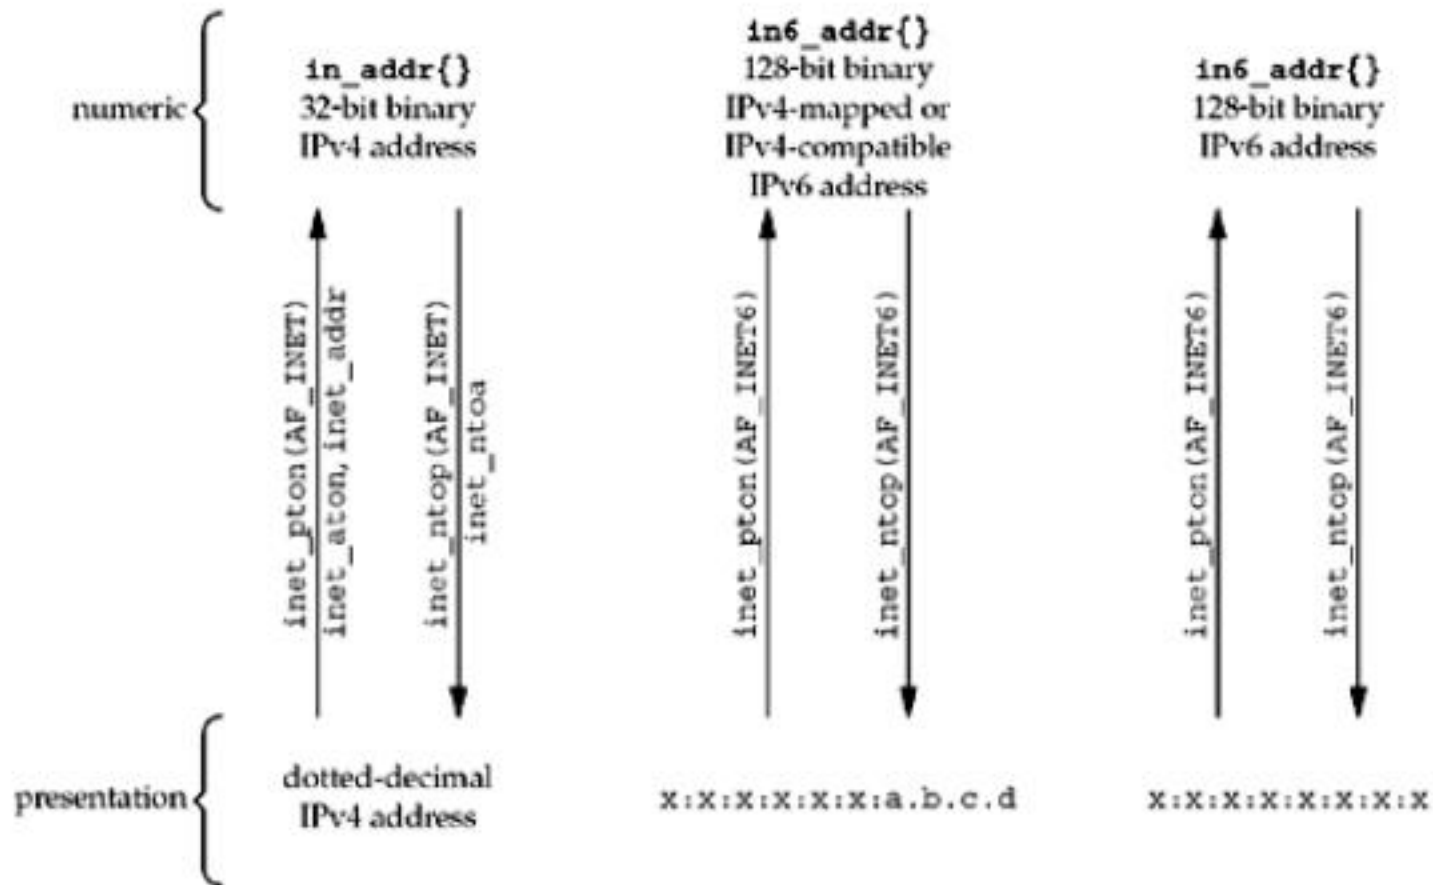
Returns: 1 if OK, 0 if input not a valid presentation format, -1 on error

```
const char *inet_ntop(int family, const void *addrptr, char *strptr, size_t len);
```

Returns: pointer to result if OK, NULL on error

- The *family* argument for both functions is either AF_INET or AF_INET6

Summary of address conversion functions



Elementary TCP Sockets

Socket functions for elementary TCP client/server

